

Data Preprocessing:

Load dataset and split data into training, validation, and testing data

```
import sklearn
import numpy as np
from operator import index
from numpy import number
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import seaborn as sn
from sklearn.metrics import classification_report, confusion_matrix

# Load breast cancer dataset
data = sklearn.datasets.load_breast_cancer()

label = data.target
data = data.data

# Divide data into training, validation, and testing data. Data is
# trainingFeatures, remainingFeatures, trainingLabels, remainingLabel
# validationFeatures, testingFeatures, validationLabels, testingLabel
```

Creating a "Node" class, which creates the nodes found in the tree

```
class Node:
    def __init__(self, featureChosen=None, threshold=None, left=None,
                 # Intializing each node with the values provided to it from the
                 # Each node would contain the information gain associated with
                 self.featureChosen = featureChosen
                 self.threshold = threshold
                 self.left = left
                 self.right = right
                 self.informationGain = informationGain
                 self.entropy = entropy
                 self.value = value
```

Creating a "Tree" class, which handles the logic for the implementation of the entire tree

```
class DecisionTree:
    def __init__(self, minSamplesSplit=2, maxDepth=2):
```

```
def __init__(self, minSamplesSplit=2, maxDepth=2):
    # Initializes the root of the tree. Placed as None as the tree or:
    self.root = None
    # Identifies the stopping condition for the tree.
    self.minSamplesSplit = minSamplesSplit
    self.maxDepth = maxDepth

def makeTree(self, trainingFeatures, trainingLabels, currentDepth=0):
    numberofSamples = trainingLabels.shape[0]
    numberofFeatures = trainingFeatures.shape[1]
    pureSplit = len(np.unique(trainingLabels)) == 1

    # Account for stopping conditions
    if currentDepth <= self.maxDepth and numberofSamples >= self.minSamplesSplit:
        # Get optimal split
        optimalSplit = self.getOptimalSplit(trainingFeatures, trainingLabels)

        # Ensure current optimal split has a positive information gain
        if optimalSplit["informationGain"] > 0:
            # Build left subtree
            leftSubtree = self.makeTree(optimalSplit["leftFeatures"], optimalSplit["leftLabels"], currentDepth + 1)

            # Build right subtree
            rightSubtree = self.makeTree(optimalSplit["rightFeatures"], optimalSplit["rightLabels"], currentDepth + 1)

            return Node(optimalSplit["featureChosen"], optimalSplit["threshold"], leftSubtree, rightSubtree)

    # If stopping conditions are met, then this is currently a leaf node
    leaf = self.makeLeafNode(trainingLabels)
    return Node(value=leaf)

def getOptimalSplit(self, trainingFeatures, trainingLabels, numberofFeatures):
    # Create dictionary to store best split
    bestSplit = {}
    # Initialize the maximum information gain with minus infinity
    # This is crucial as the greedy algorithm must initially start from negative infinity
    maxInformationGain = -float("inf")

    # Loop over all possible features to determine which feature will yield the highest information gain
    for featureIndex in range(numberofFeatures):
        featureValues = trainingFeatures[:, featureIndex]
        possibleThresholds = np.unique(featureValues)
        for threshold in possibleThresholds:
            # Split according to current given values
            dataLeft, dataRight, labelsLeft, labelsRight = self.split(trainingFeatures, threshold)
            # Ensure both left and right children are not null
            if len(dataLeft) > 0 and len(dataRight) > 0:
                informationGain = self.informationGain(trainingLabels, labelsLeft, labelsRight)
                # Update current split if the information gain obtained from this split is higher than the previous one
                if informationGain > maxInformationGain:
                    bestSplit["featureChosen"] = featureIndex
                    bestSplit["threshold"] = threshold
                    maxInformationGain = informationGain
```

```
        if informationGain > maxInformationGain:
            bestSplit["featureChosen"] = featureIndex
            bestSplit["threshold"] = threshold
            bestSplit["leftFeatures"] = dataLeft
            bestSplit["leftLabels"] = labelsLeft
            bestSplit["rightFeatures"] = dataRight
            bestSplit["rightLabels"] = labelsRight
            bestSplit["informationGain"] = informationGain
            maxInformationGain = informationGain

    # Return the current best split
    return bestSplit

def split(self, trainingFeatures, trainingLabels, featureIndex, th):
    # Split the data and labels into left and right data and features
    dataLeft = np.array([row for row in trainingFeatures if row[featureIndex] <= threshold])
    dataRight = np.array([row for row in trainingFeatures if row[featureIndex] > threshold])
    labelsLeft = trainingLabels[trainingFeatures[:, featureIndex] <= threshold]
    labelsRight = trainingLabels[trainingFeatures[:, featureIndex] > threshold]
    return dataLeft, dataRight, labelsLeft, labelsRight

def makeLeafNode(self, trainingLabels):
    #
    trainingLabels = np.array(trainingLabels)
    return max(set(trainingLabels), key=list(trainingLabels).count)

def getEntropy(self, trainingLabels):
    labels = np.unique(trainingLabels)
    # Initialize entropy with zero
    entropy = 0
    for label in labels:
        # Calculate the probability of occurrence of the label
        probability = len(trainingLabels[trainingLabels == label]) / len(trainingLabels)
        entropy += -probability * np.log2(probability)
    return entropy

def informationGain(self, labelCurrent, labelLeft, labelRight):
    # Get the weight of the left and right children from the total number of samples
    weightedLeft = len(labelLeft) / len(labelCurrent)
    weightedRight = len(labelRight) / len(labelCurrent)

    # Calculate the information gain
    infoGain = self.getEntropy(labelCurrent) - (weightedLeft * self.getEntropy(labelLeft) + weightedRight * self.getEntropy(labelRight))
    return infoGain

def printTree(self, tree=None, indent=" "):
    if not tree:
        tree = self.root
```

```
if tree.value is not None:  
    if tree.value == 0:  
        print("Malignant")  
    elif tree.value == 1:  
        print("Benign")  
    else:  
        print(tree.value)  
else:  
    print("X"+str(tree.featureChosen), "<=", tree.threshold, "?")  
    print("%sInformation Gain: " % (indent),tree.informationGain)  
    print("%sleft:" % (indent), end="")  
    self.printTree(tree.left, indent + indent)  
    print("%sright:" % (indent), end="")  
    self.printTree(tree.right, indent + indent)  
  
def fit(self, trainingFeatures, trainingLabels):  
    self.root = self.makeTree(trainingFeatures, trainingLabels)  
  
def predict(self, testingFeatures):  
    predictions = [self.makePredictions(row, self.root) for row in testingFeatures]  
    return predictions  
  
def makePredictions(self, testingFeatures, tree):  
    # If currently at a leaf node, predict  
    if tree.value != None:  
        return tree.value  
    featureValue = testingFeatures[tree.featureChosen]  
    # Else compare the featureValue to the current level tree threshold  
    if featureValue <= tree.threshold:  
        return self.makePredictions(testingFeatures, tree.left)  
    else:  
        # Else, the right subtree will be traversed  
        return self.makePredictions(testingFeatures, tree.right)  
  
def accuracy(self, testingFeatures, testingLabels):  
    predictions = self.predict(testingFeatures)  
    return np.mean(np.array(predictions) == np.array(testingLabels))
```

```
tree = DecisionTree(minSamplesSplit=3,maxDepth=3)
tree.fit(trainingFeatures, trainingLabels)
tree.printTree()

X22 <= 105.9 ?
Information Gain:  0.645947668978021
left:X27 <= 0.1423 ?
Information Gain:  0.09557508006087603
left:X3 <= 690.2 ?
Information Gain:  0.03145500992451482
left:X14 <= 0.00328 ?
Information Gain:  0.023338451601787248
left:Benign
right:Benign
right:Malignant
right:X13 <= 18.15 ?
Information Gain:  0.6394725269414906
left:X1 <= 20.22 ?
Information Gain:  0.6500224216483541
left:Benign
right:Malignant
right:Malignant
right:X27 <= 0.1489 ?
Information Gain:  0.2080739414679902
left:X21 <= 19.31 ?
Information Gain:  0.29469601591956374
left:Benign
right:X20 <= 16.76 ?
Information Gain:  0.3515358797217579
left:Benign
right:Malignant
right:Malignant
```

```
# Testing the tree
print("Tree Accuracy on Validation Data:", tree.accuracy(validation
```

```
Tree Accuracy on Validation Data: 0.9176470588235294
```

Fine Tuning the Hyperparameters:

Maximum Tree Depth and Minimum Samples Per Split

```
def tuning(minSamplesSplit, maxDepth):
    # Create a new instance of a tree with every pass
    trainingTree = DecisionTree(minSamplesSplit=minSamplesSplit, maxDe
    # Fit the tree with the training labels and features
    trainingTree.fit(trainingFeatures, trainingLabels)
    # Test the tree with the validation labels and features
    return trainingTree.accuracy(validationFeatures, validationLabels)

def multipleRuns(minSamplesSplit=[2,5,10], maxDepth=[2,4,6,8,10]):
    # Sentinel values (initial guard values)
    optimalSampleSplit = 0
    optimalMaxDepth = 0
    maxAccuracy = 0

    # Attempting every possible combination of minimum samples split
    for sampleSplit in minSamplesSplit:
        for depth in maxDepth:
            accuracy = tuning(sampleSplit, depth)
            # Changing the values of the maximum accuracy, optimal sample
            if accuracy > maxAccuracy:
                maxAccuracy = accuracy
                optimalSampleSplit = sampleSplit
                optimalMaxDepth = depth
    print("Hyperparameter Tuning:")
    print("Minimum Samples Per Split:", optimalSampleSplit)
    print("Maximum Tree Depth:", optimalMaxDepth)
    print("Maximum Accuracy:", maxAccuracy)
    return optimalSampleSplit, optimalMaxDepth, maxAccuracy

minSampleSplit, maxDepth, maxAccuracy = multipleRuns()
```

```
Hyperparameter Tuning:
Minimum Samples Per Split: 2
Maximum Tree Depth: 4
Maximum Accuracy: 0.8941176470588236
```

```
tree = DecisionTree(minSamplesSplit=minSampleSplit, maxDepth=maxDept
tree.fit(trainingFeatures, trainingLabels)
# Test optimal minSamplesSplit and maxDepth on testing data
predictions = tree.predict(testingFeatures)
print("Tree Accuracy on Testing Data:", tree.accuracy(testingFeatur
```

```
Tree Accuracy on Testing Data: 0.9302325581395349
```

```
# Examining training and validation accuracy change with maximum de
samples=[2,4,6,8,10]
```

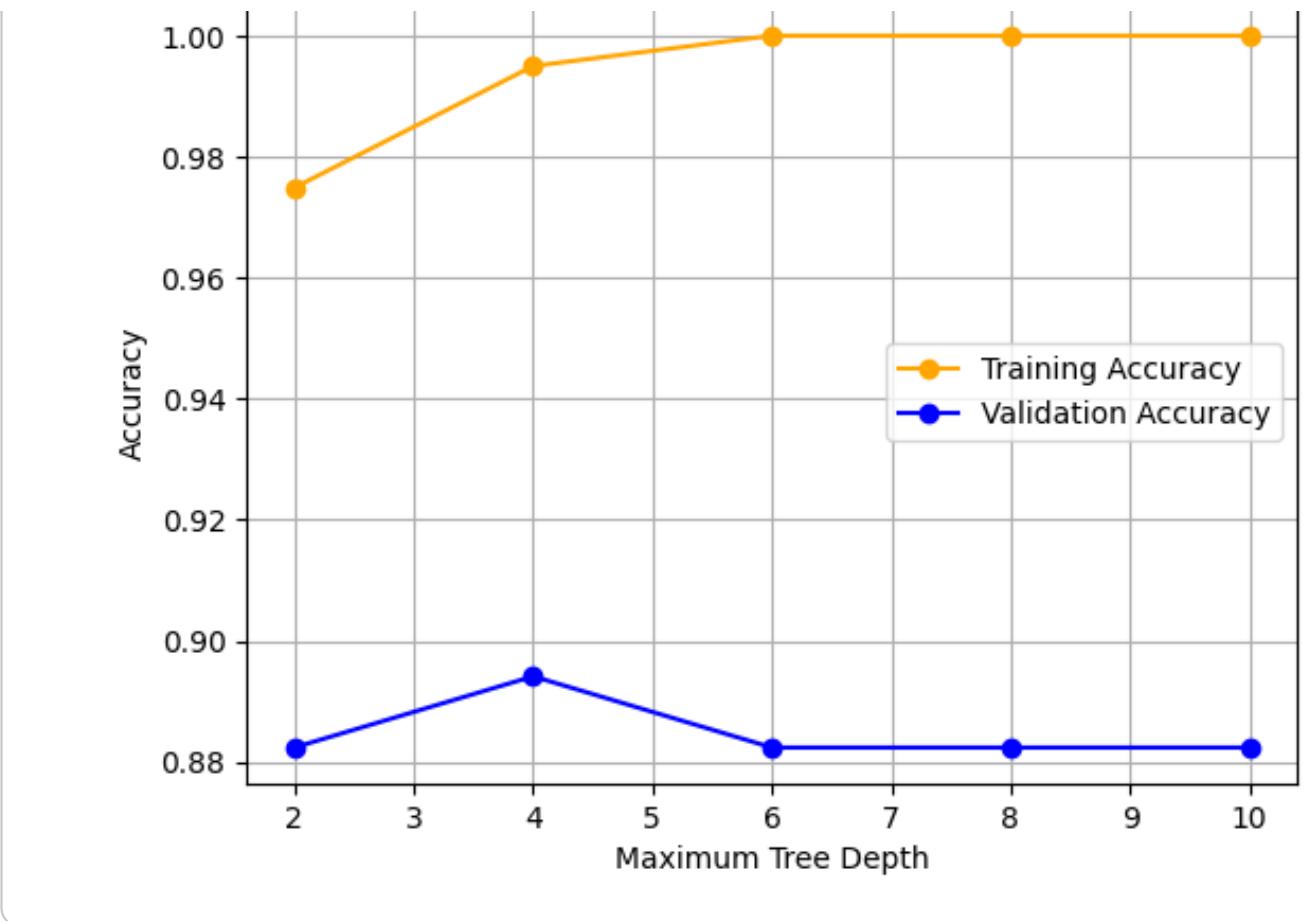
```
trainingAccuracyData=[]
validationAccuracyData=[]
def accuracyChange(samples=[2,4,6,8,10]):
    # Changing the maximum depth only while keeping the minimum sample size
    for depth in samples:
        # Create a new instance of a tree with every possible feature
        trainingTree = DecisionTree(minSamplesSplit=2,maxDepth=depth)
        trainingTree.fit(trainingFeatures, trainingLabels)
        # Store the tree's accuracy on the training and validation data
        trainingAccuracy = trainingTree.accuracy(trainingFeatures, trainingLabels)
        validationAccuracy = trainingTree.accuracy(validationFeatures, validationLabels)
        print("Maximum Tree Depth:", depth)
        print("Training Accuracy:", trainingAccuracy)
        print("Validation Accuracy:", validationAccuracy)
        trainingAccuracyData.append(trainingAccuracy)
        validationAccuracyData.append(validationAccuracy)

# Plotting
def plotAccuracyChange(samples, trainingAccuracyData, validationAccuracyData):
    plt.plot(samples, trainingAccuracyData, label='Training Accuracy')
    plt.plot(samples, validationAccuracyData, label='Validation Accuracy')
    plt.xlabel('Maximum Tree Depth')
    plt.ylabel('Accuracy')
    plt.legend()
    plt.grid()
    plt.title('Maximum Tree Depth VS Accuracy')
    plt.show()

accuracyChange()
plotAccuracyChange(samples, trainingAccuracyData, validationAccuracyData)
```

```
Maximum Tree Depth: 2
Training Accuracy: 0.9748743718592965
Validation Accuracy: 0.8823529411764706
Maximum Tree Depth: 4
Training Accuracy: 0.9949748743718593
Validation Accuracy: 0.8941176470588236
Maximum Tree Depth: 6
Training Accuracy: 1.0
Validation Accuracy: 0.8823529411764706
Maximum Tree Depth: 8
Training Accuracy: 1.0
Validation Accuracy: 0.8823529411764706
Maximum Tree Depth: 10
Training Accuracy: 1.0
Validation Accuracy: 0.8823529411764706
```

Maximum Tree Depth VS Accuracy



The results of the accuracy of the training and validation data versus the tree depth present the following findings:

- 1- The deeper the tree was allowed to become, the easier it overfitted.
- 2- A depth of 6 or more would leave the tree having well defined aspects for the training data (no misses would result and the accuracy of the training data would be 100%)
- 3- The performance of the validation data does not vary much with the change in the depth. This is predicted as the tree originally splits on the most distinctive features for 70% of the provided data (training data). A similar trend would be expected in the validation and testing data.

```
tree = DecisionTree(minSamplesSplit=minSampleSplit,maxDepth=maxDept  
tree.fit(trainingFeatures, trainingLabels)  
# Test optimal minSamplesSplit and maxDepth on testing data  
predictions = tree.predict(testingFeatures)  
  
numpyPredictions = np.array(predictions)  
numpyActualLabels = np.array(testingLabels)
```

```
# Performance Metrics

# Calculating the precision, recall, and f-score
# Precision: How many positively predicted instances the model pred
# Recall: Out of all of the positive instances, how many did the mo
# F1-Score: Harmonic mean of both the precision and the recall. Hea
true_positives = np.sum(np.logical_and(numpyPredictions == 0, numpy
true_negatives = np.sum(np.logical_and(numpyPredictions == 1, numpy
false_positives = np.sum(np.logical_and(numpyPredictions == 0, numpy
false_negatives = np.sum(np.logical_and(numpyPredictions == 1, numpy

precision = true_positives / (true_positives + false_positives)
recall = true_positives / (true_positives + false_negatives)
f1_score = 2 * ((precision * recall) / (precision + recall))

print("Classification Report for Testing Data:")
print("-----")
print("Benign: Label 0 || Malignant: Label 1")
print("Precision: ",precision)
print("Recall: ",recall)
print("F1-Score: ",f1_score)
print("-----")

# Confusion Matrix

print("Confusion Matrix for Testing Data:")
print("-----")

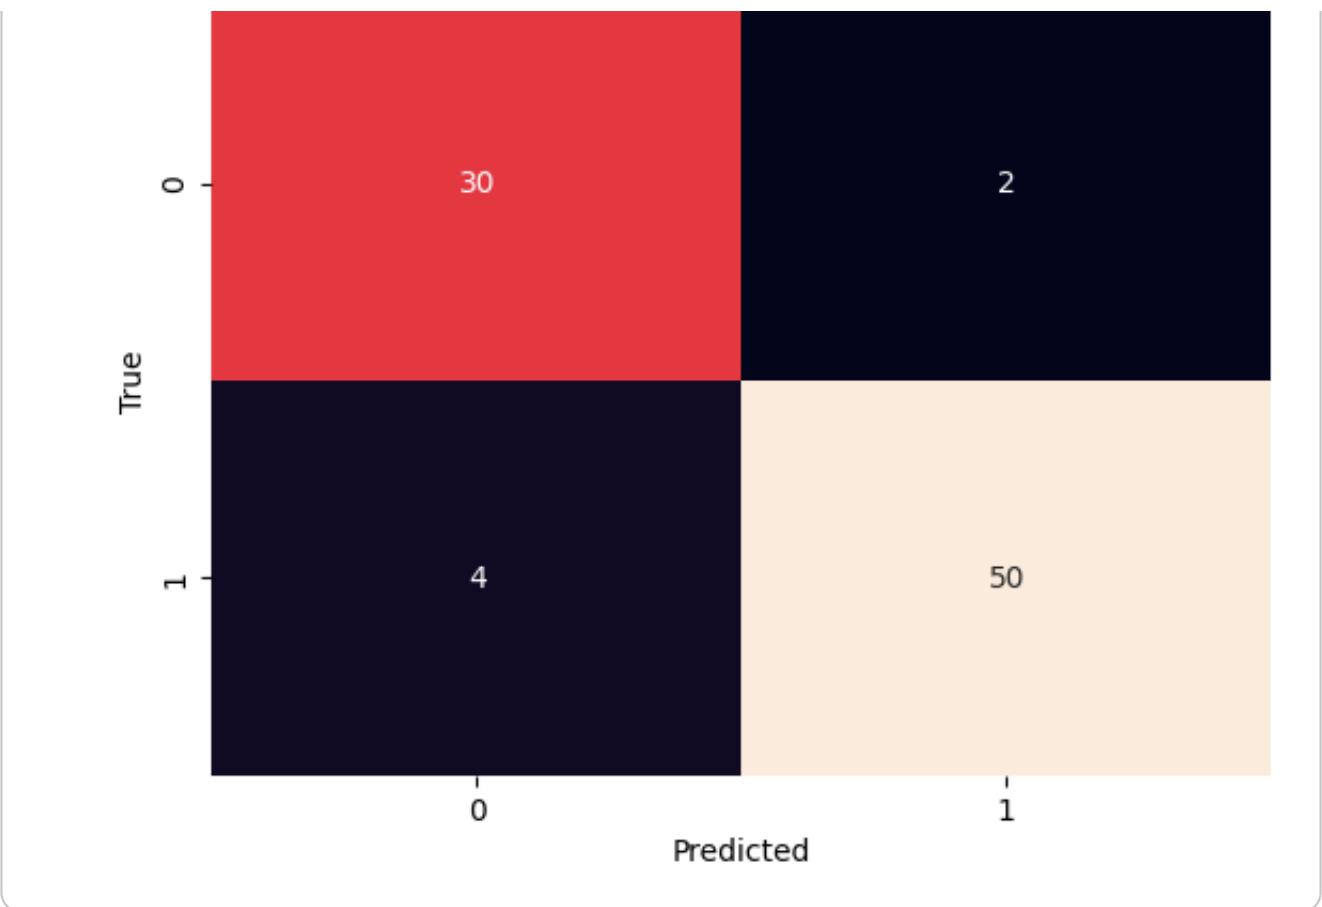
confusionMatrix = confusion_matrix(testingLabels, predictions)
plt.figure()
sn.heatmap(confusionMatrix, annot=True, fmt="d", cbar=False)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix for Testing Data:')
plt.show()
```

Classification Report for Testing Data:

Benign: Label 0 || Malignant: Label 1
Precision: 0.8823529411764706
Recall: 0.9375
F1-Score: 0.9090909090909091

Confusion Matrix for Testing Data:

Confusion Matrix for Testing Data:

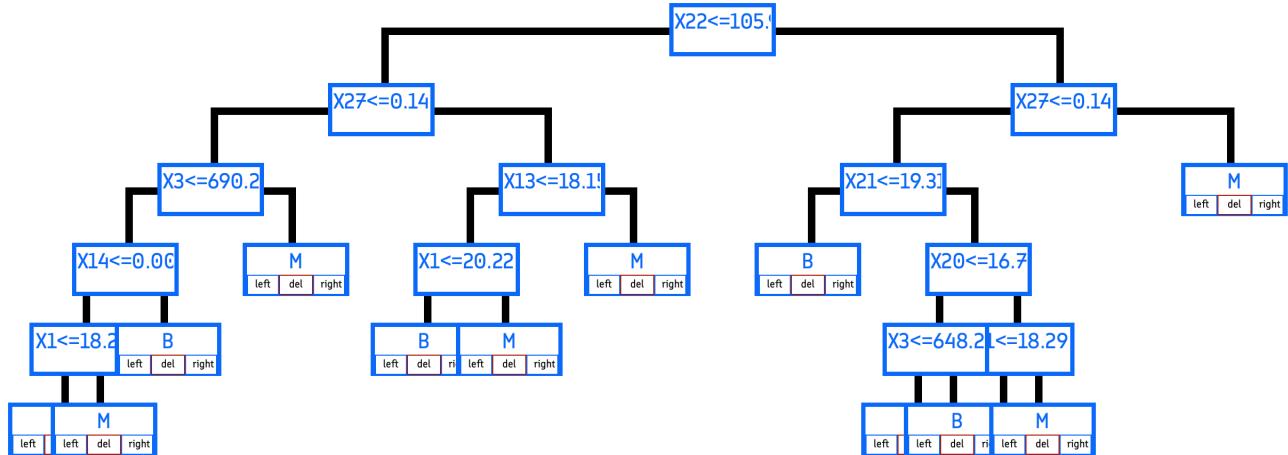


The tree correctly identifies 30 malignant cases and 50 benign cases. Unfortunately, the tree misses out on 6 samples, 4 of which are benign but predicted malignant and 2 of which are malignant but predicted benign. Although the last two samples are definitely misleading, the tree's main false identifications are tolerable ones (classifying as malignant while being benign, which is much better than the other way around).

```
# Commenting on the feature importance by printing the tree:  
tree.printTree()  
  
X22 <= 105.9 ?  
    Information Gain:  0.645947668978021  
    left:X27 <= 0.1423 ?  
        Information Gain:  0.09557508006087603  
        left:X3 <= 690.2 ?  
            Information Gain:  0.03145500992451482  
            left:X14 <= 0.00328 ?  
                Information Gain:  0.023338451601787248  
                left:X1 <= 18.22 ?  
                    Information Gain:  0.6500224216483541  
                    left:Benign  
                    right:Malignant  
                right:Benign  
            right:Malignant  
        right:X13 <= 18.15 ?  
            Information Gain:  0.6394725269414906  
            left:X1 <= 20.22 ?  
                Information Gain:  0.6500224216483541  
                left:Benign  
                right:Malignant  
            right:Malignant  
        right:X27 <= 0.1489 ?  
            Information Gain:  0.2080739414679902  
            left:X21 <= 19.31 ?  
                Information Gain:  0.29469601591956374  
                left:Benign  
                right:X20 <= 16.76 ?  
                    Information Gain:  0.3515358797217579  
                    left:X3 <= 648.2 ?  
                        Information Gain:  0.5567796494470396  
                        left:Malignant  
                        right:Benign  
                    right:X1 <= 18.29 ?  
                        Information Gain:  0.13509531357637425  
                        left:Malignant  
                        right:Malignant  
                right:Malignant
```

The output of the tree print portrays how every split was done to determine the most distinctive features chosen to split on. Each split was made according to the feature split that yielded the highest information gain.

Tree Visualized



Given the provided image, the most distinctive feature that initially splits the tree is X22, followed by X27. These two are shared by both trees in both levels.