

## Data Preprocessing:

Load dataset and split data into training, validation, and testing data

```
import sklearn
import numpy as np
from operator import index
from numpy import number
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import seaborn as sn
from sklearn.metrics import classification_report, confusion_matrix

# Load breast cancer dataset
data = sklearn.datasets.load_breast_cancer()

label = data.target
data = data.data

# Divide data into training, validation, and testing data. Data is divided into
# trainingFeatures, remainingFeatures, trainingLabels, remainingLabels
# validationFeatures, testingFeatures, validationLabels, testingLabels
```

Creating a "Node" class, which creates the nodes found in the tree

```
class Node:
    def __init__(self, featureChosen=None, threshold=None, left=None,
                 # Initializing each node with the values provided to it from the
                 # Each node would contain the information gain associated with
                 self.featureChosen = featureChosen
                 self.threshold = threshold
                 self.left = left
                 self.right = right
                 self.informationGain = informationGain
                 self.entropy = entropy
                 self.value = value
```

Creating a "Tree" class, which handles the logic for the implementation of the entire tree

```
class DecisionTree:
```

```
def __init__(self, minSamplesSplit=2, maxDepth=2):
    # Initializes the root of the tree. Placed as None as the tree is initially empty.
    self.root = None
    # Identifies the stopping condition for the tree.
    self.minSamplesSplit = minSamplesSplit
    self.maxDepth = maxDepth

def makeTree(self, trainingFeatures, trainingLabels, currentDepth=0):
    numberofSamples = trainingLabels.shape[0]
    numberofFeatures = trainingFeatures.shape[1]
    pureSplit = len(np.unique(trainingLabels)) == 1

    # Account for stopping conditions
    if currentDepth <= self.maxDepth and numberofSamples >= self.minSamplesSplit:
        # Get optimal split
        optimalSplit = self.getOptimalSplit(trainingFeatures, trainingLabels)

        # Ensure current optimal split has a positive information gain
        if optimalSplit["informationGain"] > 0:
            # Build left subtree
            leftSubtree = self.makeTree(optimalSplit["leftFeatures"], trainingLabels)

            # Build right subtree
            rightSubtree = self.makeTree(optimalSplit["rightFeatures"], trainingLabels)

            return Node(optimalSplit["featureChosen"], optimalSplit["threshold"], leftSubtree, rightSubtree)

    # If stopping conditions are met, then this is currently a leaf
    leaf = self.makeLeafNode(trainingLabels)
    return Node(value=leaf)

def getOptimalSplit(self, trainingFeatures, trainingLabels, numberofFeatures):
    # Create dictionary to store best split
    bestSplit = {}
    # Initialize the maximum information gain with minus infinity
    # This is crucial as the greedy algorithm must initially start from a large value
    maxInformationGain = -float("inf")

    # Loop over all possible features to determine which feature will provide the best split
    for featureIndex in range(numberofFeatures):
        featureValues = trainingFeatures[:, featureIndex]
        possibleThresholds = np.unique(featureValues)
        for threshold in possibleThresholds:
            # Split according to current given values
            dataLeft, dataRight, labelsLeft, labelsRight = self.split(featureIndex, threshold)
            # Ensure both left and right children are not null
            if len(dataLeft) > 0 and len(dataRight) > 0:
                informationGain = self.informationGain(trainingLabels, labelsLeft, labelsRight)
                if informationGain > maxInformationGain:
                    bestSplit[featureIndex] = threshold
                    maxInformationGain = informationGain

    return bestSplit
```

```
# Update current split if the information gain obtained f
if informationGain > maxInformationGain:
    bestSplit["featureChosen"] = featureIndex
    bestSplit["threshold"] = threshold
    bestSplit["leftFeatures"] = dataLeft
    bestSplit["leftLabels"] = labelsLeft
    bestSplit["rightFeatures"] = dataRight
    bestSplit["rightLabels"] = labelsRight
    bestSplit["informationGain"] = informationGain
    maxInformationGain = informationGain

# Return the current best split
return bestSplit

def split(self, trainingFeatures, trainingLabels, featureIndex, t
dataLeft = np.array([row for row in trainingFeatures if row[fea
dataRight = np.array([row for row in trainingFeatures if row[fe
labelsLeft = trainingLabels[trainingFeatures[:, featureIndex] <
labelsRight = trainingLabels[trainingFeatures[:, featureIndex]
return dataLeft, dataRight, labelsLeft, labelsRight

def makeLeafNode(self, trainingLabels):
    trainingLabels = np.array(trainingLabels)
    return max(set(trainingLabels), key=list(trainingLabels).count)

def getEntropy(self, trainingLabels):
    labels = np.unique(trainingLabels)
    # Initialize entropy with zero
    entropy = 0
    for label in labels:
        # Calculate the probability of occurrence of the label
        probability = len(trainingLabels[trainingLabels == label]) /
        entropy += -probability * np.log2(probability)
    return entropy

def informationGain(self, labelCurrent, labelLeft, labelRight):
    # Get the weight of the left and right children from the total
    weightedLeft = len(labelLeft) / len(labelCurrent)
    weightedRight = len(labelRight) / len(labelCurrent)

    # Calculate the information gain
    infoGain = self.getEntropy(labelCurrent) - (weightedLeft * self
    return infoGain

def printTree(self, tree=None, indent=" "):
    if not tree:
        tree = self.root
    if tree.value is not None:
```

```
if tree.value == 1:
    print("Malignant")
elif tree.value == 0:
    print("Benign")
else:
    print(tree.value)
else:
    print("X"+str(tree.featureChosen), "<=", tree.threshold, "?")
    print("%sInformation Gain: " % (indent), tree.informationGain)
    print("%sleft:" % (indent), end="")
    self.printTree(tree.left, indent + indent)
    print("%sright:" % (indent), end="")
    self.printTree(tree.right, indent + indent)

def fit(self, trainingFeatures, trainingLabels):
    self.root = self.makeTree(trainingFeatures, trainingLabels)

def predict(self, testingFeatures):
    predictions = [self.makePredictions(row, self.root) for row in
    return predictions

def makePredictions(self, testingFeatures, tree):
    if tree.value != None:
        return tree.value
    featureValue = testingFeatures[tree.featureChosen]
    if featureValue <= tree.threshold:
        return self.makePredictions(testingFeatures, tree.left)
    else:
        return self.makePredictions(testingFeatures, tree.right)

def accuracy(self, testingFeatures, testingLabels):
    predictions = self.predict(testingFeatures)
    return np.mean(np.array(predictions) == np.array(testingLabels))
```

```
tree = DecisionTree(minSamplesSplit=3,maxDepth=3)
tree.fit(trainingFeatures, trainingLabels)
tree.printTree()

X22 <= 105.9 ?
Information Gain:  0.645947668978021
left:X27 <= 0.1423 ?
Information Gain:  0.09557508006087603
left:X3 <= 690.2 ?
Information Gain:  0.03145500992451482
left:X14 <= 0.00328 ?
Information Gain:  0.023338451601787248
left:Malignant
right:Malignant
right:Benign
right:X13 <= 18.15 ?
Information Gain:  0.6394725269414906
left:X1 <= 20.22 ?
Information Gain:  0.6500224216483541
left:Malignant
right:Benign
right:Benign
right:X27 <= 0.1489 ?
Information Gain:  0.2080739414679902
left:X21 <= 19.31 ?
Information Gain:  0.29469601591956374
left:Malignant
right:X20 <= 16.76 ?
Information Gain:  0.3515358797217579
left:Malignant
right:Benign
right:Benign
```

```
# Testing the tree
print("Tree Accuracy on Validation Data:", tree.accuracy(validation
```

```
Tree Accuracy on Validation Data: 0.9176470588235294
```

Fine Tuning the Hyperparameters:

Maximum Tree Depth and Minimum Samples Per Split

```

def tuning(minSamplesSplit, maxDepth):
    trainingTree = DecisionTree(minSamplesSplit=minSamplesSplit,maxDe
    trainingTree.fit(trainingFeatures, trainingLabels)
    return trainingTree.accuracy(validationFeatures, validationLabels)

def multipleRuns(minSamplesSplit=[2,5,10], maxDepth=[2,4,6,8,10]):
    optimalSampleSplit = 0
    optimalMaxDepth = 0
    maxAccuracy = 0

    for sampleSplit in minSamplesSplit:
        for depth in maxDepth:
            accuracy = tuning(sampleSplit, depth)
            if accuracy > maxAccuracy:
                maxAccuracy = accuracy
                optimalSampleSplit = sampleSplit
                optimalMaxDepth = depth
    print("Hyperparameter Tuning:")
    print("Minimum Samples Per Split:", optimalSampleSplit)
    print("Maximum Tree Depth:", optimalMaxDepth)
    print("Maximum Accuracy:", maxAccuracy)
    return optimalSampleSplit, optimalMaxDepth, maxAccuracy

minSampleSplit, maxDepth, maxAccuracy = multipleRuns()

```

Hyperparameter Tuning:  
 Minimum Samples Per Split: 2  
 Maximum Tree Depth: 4  
 Maximum Accuracy: 0.8941176470588236

```

tree = DecisionTree(minSamplesSplit=minSampleSplit,maxDepth=maxDept
tree.fit(trainingFeatures, trainingLabels)
# Test optimal minSamplesSplit and maxDepth on testing data
predictions = tree.predict(testingFeatures)
print("Tree Accuracy on Testing Data:", tree.accuracy(testingFeatur

```

Tree Accuracy on Testing Data: 0.9302325581395349

```

# Examining training and validation accuracy change with maximum de
samples=[2,4,6,8,10]
trainingAccuracyData=[]
validationAccuracyData=[]
def accuracyChange(samples=[2,4,6,8,10]):
    for depth in samples:
        trainingTree = DecisionTree(minSamplesSplit=2,maxDepth=depth)
        trainingTree.fit(trainingFeatures, trainingLabels)


```


```
trainingAccuracy = trainingTree.accuracy(trainingFeatures, train
validationAccuracy = trainingTree.accuracy(validationFeatures,
print("Maximum Tree Depth:", depth)
print("Training Accuracy:", trainingAccuracy)
print("Validation Accuracy:", validationAccuracy)
trainingAccuracyData.append(trainingAccuracy)
validationAccuracyData.append(validationAccuracy)

def plotAccuracyChange(samples, trainingAccuracyData, validationAcc
plt.plot(samples, trainingAccuracyData, label='Training Accuracy'
plt.plot(samples, validationAccuracyData, label='Validation Accur
plt.xlabel('Maximum Tree Depth')
plt.ylabel('Accuracy')
plt.legend()
plt.grid()
plt.title('Maximum Tree Depth VS Accuracy')
plt.show()

accuracyChange()
plotAccuracyChange(samples, trainingAccuracyData, validationAccurac
```

```
Maximum Tree Depth: 2
Training Accuracy: 0.9748743718592965
Validation Accuracy: 0.8823529411764706
Maximum Tree Depth: 4
Training Accuracy: 0.9949748743718593
Validation Accuracy: 0.8941176470588236
Maximum Tree Depth: 6
Training Accuracy: 1.0
Validation Accuracy: 0.8823529411764706
Maximum Tree Depth: 8
Training Accuracy: 1.0
Validation Accuracy: 0.8823529411764706
Maximum Tree Depth: 10
Training Accuracy: 1.0
Validation Accuracy: 0.8823529411764706
```





```
tree = DecisionTree(minSamplesSplit=minSampleSplit,maxDepth=maxDepth)
tree.fit(trainingFeatures, trainingLabels)
# Test optimal minSamplesSplit and maxDepth on testing data
predictions = tree.predict(testingFeatures)

numpyPredictions = np.array(predictions)
numpyActualLabels = np.array(testingLabels)

# Performance Metrics

# Calculating the precision, recall, and f-score
# Precision: How many positively predicted instances the model predicted
# Recall: Out of all of the positive instances, how many did the model predict
# F1-Score: Harmonic mean of both the precision and the recall. Heavily weighted by precision
true_positives = np.sum(np.logical_and(numpyPredictions == 0, numpyActualLabels == 0))
true_negatives = np.sum(np.logical_and(numpyPredictions == 1, numpyActualLabels == 1))
false_positives = np.sum(np.logical_and(numpyPredictions == 0, numpyActualLabels == 1))
false_negatives = np.sum(np.logical_and(numpyPredictions == 1, numpyActualLabels == 0))

precision = true_positives / (true_positives + false_positives)
recall = true_positives / (true_positives + false_negatives)
f1_score = 2 * ((precision * recall) / (precision + recall))

print("Classification Report for Testing Data:")
print("-----")
print("Benign: Positive || Malignant: Negative")
print("Precision: ",precision)
print("Recall: ",recall)
print("F1-Score: ",f1_score)
print("-----")

# Confusion Matrix

print("Confusion Matrix for Testing Data:")
print("-----")
```

```
confusionMatrix = confusion_matrix(testingLabels, predictions)
plt.figure()
sn.heatmap(confusionMatrix, annot=True, fmt="d", cbar=False)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix for Testing Data:')
plt.show()
```

Classification Report for Testing Data:

-----  
Benign: Positive || Malignant: Negative  
Precision: 0.8823529411764706  
Recall: 0.9375  
F1-Score: 0.9090909090909091

-----  
Confusion Matrix for Testing Data:

